





Creating architecture.md file for resume retrieval using RAG approach:





My Workspace

Search

Invite Upgrade

GenAI training

Step by step resume retrieval using RAG approach

POST Step 2: Generate 1024D vector embedding for single text string using Mistral

POST http://localhost:3000/v1/embeddings

Send

Docs Params Authorization Headers 10 Body Scripts Tests Settings

Query Params

Key	Value	Description	Bulk Edit
Key	Value	Description	

Body Cookies Headers 10 Test Results

200 OK - 1.94 s - 19.23 KB

Save Response

JSON

```
{
  "object": "list",
  "data": [
    {
      "object": "embedding",
      "index": 0,
      "embedding": [
        -0.66389494296875,
        0.0261077880859375,
        0.0458984375,
        -0.0495666546875,
        0.04107666015625,
        0.01654852734375,
        -0.00016645678373535156,
        0.0012731552124023438,
        0.0287628173828125,
        -0.0157928466796875,
        ...
      ]
    }
  ]
}
```

Activate Windows
Go to Settings to activate Windows.

Connect Git Console Terminal

23:24 12-06-2026

File Edit Selection View Go Run Terminal Help

Untitled (Workspace) - Antigravity IDE - README.md

Explorer

- Untitled (Workspace)
- scratch
- dist
- node_modules
- prompts
- rag_ingestion_resume
- rag_resume_ingestion
- src
- uploads
- .env
- .gitignore
- package-lock.json
- package.json
- README.md
- Resume ingestion guide.docx
- resume_ingestion_architecture.md
- Resume_Ingestion_Steps.docx
- ResumeIngestion_Prompt.txt
- Step by step resume ingestion using RAG ...
- tsconfig.json
- rag_resume_retrieval
- copilot_prompts
- node_modules
- config
- mongojs
- controller

Backend Service > API Endpoints Guide > Match Summarization

```
1 ## Resume Search Platform: Backend
56 ## API Endpoints Guide
337 ### 7. Match Summarization API (Match
382 ---
383
384 ### 8. End-to-End Resume Search
385 Pipeline API
386 * **Endpoints**:
387 - 'GET /v1/search'
388 - 'GET /api/v1/search'
389 * **Query Parameters**:
390 - 'q' (string, required): The search
391 text/phrase key.
392 * **Use Case**: Runs RRF hybrid
393 retrieval (combining Vector Search and
394 BM25 Search) to collect relevant
395 candidate profiles, slices the top 5
396 candidates, triggers a batch LLM
397 explanation summary request using the
398 Groq API ('meta-llama/
399
400 and-08-01-7fkhbk.mongodb.net:27017,ac-9nj1sya-shar
401 d-08-02-7fkhbk.mongodb.net:27017/resumedb?ssl=true
402 &replicaSet=atlas-2a9fa6-shard-0&authSource=admin&r
403 etryWrites=true&majority
404 dbName='resumedb'
405 [2026-06-12 23:17:14.965 +0530] INFO: Server is run
406 ning in development mode on port 3000
407 [2026-06-12 23:17:15.407 +0530] INFO: Successfully
408 connected to MongoDB.
```

Confirming Vector Embedding Requirements

refer resume_search_architecture.md file

Generate code 'To implement Vector Embedding + '/v1/embeddings' with Architecture.md and create_endpoint.md prompt file' with mistral-embed for embeddings (store in .env) as single input at time.

Worked for 10s >

I have created the implementation plan for the Mistral-based Vector Embedding service and /v1/embeddings endpoint. Please review the plan and let me know if you approve so I can proceed with implementing and verifying the code.

Auto-proceeded with Implementation Plan

Worked for 50s >

I am waiting for the command output to verify the server status.

Implementation Plan

Implementation plan for the End-to-End Resume Search Pipeline (RRF + Batch Match

0 Files With Changes



Review Changes

Ask anything, @ to mention, / for actions

Activate Windows
Go to Settings to activate Windows.

ganeshmynampati (now) Ln 347, Col 65 (36 selected) Spaces: 2 UTF-8 LF {} Markdown Antigravity - Settings

23:19 12-06-2026



My Workspace

Search

Invited Upgrade

GenAI training

Step by step resume retrieval using RAG approach > Step 3: BM25 search

GET http://localhost:3000/v1/search/bm25?q=Java

Send

Docs Params Authorization Headers 10 Body Scripts Tests Settings Cookies

Query Params

Key	Value	Description	Bulk Edit
q	Java		
Key	Value	Description	

Body Cookies Headers 10 Test Results Save Response

200 OK 148 ms 92.41 KB

```
{
  "success": true,
  "count": 43,
  "data": [
    {
      "_id": "6a283d1974ce24f83c85ad65",
      "name": "Java Developer",
      "email": "pdineshkumar2407@gmail.com",
      "phone": "+918610849219",
      "skills": [
        "JavaScript",
        "Java",
        "SQL",
        "HTML",
        "Spring Boot",
        "jQuery",
        "Oracle",

```

Activate Windows
Go to Settings to activate Windows.

Connect Git Console Terminal 0 0 0

Globals Vault Tools

23:24
12-06-2026

File Edit Selection View Go Run Terminal Help

Untitled (Workspace) - Antigravity IDE - README.md

Explorer

- Untitled (Workspace)
- scratch
- dist
- node_modules
- prompts
- rag_ingestion_resume
- rag_resume_ingestion
- src
- uploads
- .env
- .gitignore
- package-lock.json
- package.json
- README.md
- Resume ingestion guide.docx
- resume_ingestion_architecture.md
- Resume_ingestion_steps.docx
- ResumeIngestion_Prompt.txt
- Step by step resume ingestion using RAG ...
- tsconfig.json
- rag_resume_retrieval
- copilot_prompts
- node_modules
- config
- mongojs
- controller

Backend Service > API Endpoints Guide > Match Summarization

```
1 ## Resume Search Platform: Backend
56 ## API Endpoints Guide
337 ### 7. Match Summarization API (Match
382 ---
383
384 ### 8. End-to-End Resume Search
385 Pipeline API
386 * **Endpoints**
387 - GET /v1/search
388 - GET /api/v1/search
389 * **Query Parameters**
390 - q (string, required): The search
391 text/phrase key.
392 * **Use Case**:: Runs RRF hybrid
393 retrieval (combining Vector Search and
394 BM25 Search) to collect relevant
395 candidate profiles, slices the top 5
396 candidates, triggers a batch LLM
397 explanation summary request using the
398 Groq API (meta-llama/
399
400 [2026-06-12 23:21:08.402 +0530] INFO: Executing At
401 RRF fusion completed.
402 resultsCount: 43
403 durationMs: 723
404 [2026-06-12 23:21:28.654 +0530] INFO: Calling Groq
405 API for batch candidate match explanations
406 model: "meta-llama/llama-4-scout-17b-16e-instru
407 ct"
408 count: 5
409
```

Confirming Vector Embedding Requirements

Proceed from where you left

`resume_bm25_search_index` is configured and successfully queryable.

The server is running cleanly and testing is fully complete. Please let me know if you would like to proceed with the next phase!

resume_search_architecture.md

Implement 'Vector Search config + /v1/search/vector'. DON'T use knn beta. 36 Confirm service, routes and are created and index is updated.

Worked for 16s

I have updated the implementation plan to add the Vector Search integration (GET /v1/search/vector). Please review the plan and let me know if you approve so I can proceed with the execution.

Auto-proceeded with Implementation Plan

Worked for 1m

0 Files With Changes



Review Changes



Ask anything, @ to mention, / for actions

Activate Windows
Go to Settings to activate Windows.

ganeshmynampati (now) Ln 347, Col 65 (36 selected) Spaces: 2 UTF-8 LF {} Markdown Antigravity - Settings

23:25
12-06-2026





My Workspace

Search

Invited Upgrade

GenAI training

Step by step resume retrieval using RAG approach > Step 6: Reranking

GET http://localhost:3000/v1/search/rerank?q=java

Send

Docs Params Authorization Headers 10 Body Scripts Tests Settings Cookies

Query Params

Key	Value	Description	Bulk Edit
q	java		
Key	Value	Description	

Body Cookies Headers 10 Test Results

200 OK - 2.56 s - 17.52 KB

Save Response

```
{
  "success": true,
  "count": 10,
  "data": [
    {
      "_id": "6a283d5c74ce24f83c85addc",
      "name": "RAJADURAI ARIVAZHAGAN",
      "email": "rajaduraimax96@gmail.com",
      "phone": "9690376187",
      "skills": [
        "Java",
        "Spring Boot",
        "MongoDB",
        "PostgreSQL",
        "MySQL",
        "Jenkins",
        "Git"
      ]
    }
  ]
}
```

Activate Windows
Go to Settings to activate Windows.

Connect Git Console Terminal

Globals Vault Tools

23:26 12-06-2026

File Edit Selection View Go Run Terminal Help

Untitled (Workspace) - Antigravity IDE - README.md

Confirming Vector Embedding Requirements

Proceed from where you left

- Registered GET /rerank in js search.routes.js (mounted in js app.js).

3. Verification:

- Tested GET /v1/search/rerank?q=java to verify correct re-ranking and justification returns.
- Verified 400 Bad Request handling when missing parameters.

4. Documentation:

- Documented this endpoint in README.md and walkthrough.md.

1 file changed +53-0

Review

resume_search_architecture.md

Implement 'Summarization config + /v1/search/summarize'.

11:01 PM

Worked for 16s

0 Files With Changes

Review Changes

Ask anything, @ to mention, / for actions

Activate Windows
Go to Settings to activate Windows.

Gemini 3.5 Flash (Medium)

23:20 12-06-2026

Backend Service > API Endpoints Guide > Match Summarization API

```
1 ## Resume Search Platform: Backend
56 ## API Endpoints Guide
337 ### 7. Match Summarization API (Match
382 ---
384 ### 8. End-to-End Resume Search
385 Pipeline API
386 * **Endpoints**
387 - GET /v1/search
388 - GET /api/v1/search
389 * **Query Parameters**
390 - q (string, required): The search
text/phrase key.
391 * **Use Case**:: Runs RRF hybrid
retrieval (combining Vector Search and
BM25 Search) to collect relevant
candidate profiles, slices the top 5
candidates, triggers a batch LLM
explanation summary request using the
Groq API (meta-llama/

and-08-01-7fkhbk.mongodb.net:27017,ac-9njl1sya-shar
d-08-02-7fkhbk.mongodb.net:27017/resumedb?ssl=true
&replicaSet=atlas-2o3fa6-shard-0&authSource=admin&r
etryWrites=true&majority
dbName: "resumedb"
[2026-06-12 23:17:14.965 +0530] INFO: Server is run
ning in development mode on port 3000
[2026-06-12 23:17:15.407 +0530] INFO: Successfully
connected to MongoDB.
```

node - rag_resume_retrieval

Ln 347, Col 65 (36 selected) Spaces: 2 UTF-8 LF {} Markdown Antigravity - Settings

My Workspace

Search

Invite Upgrade

GenAI training

Step by step resume retrieval using RAG approach > Step 7: Match Summarization API

GET Step 7: Match Summarization API

GET http://localhost:3000/v1/search/summarize?q=java&candidateId=6a283d5c74ce24f83c05addc

Send

Docs Params Authorization Headers 10 Body Scripts Tests Settings Cookies

Query Params

Key	Value	Description	Bulk Edit
q	java		
candidateId	6a283d5c74ce24f83c05addc		

Body Cookies Headers 10 Test Results

200 OK 817 ms 985 B Save Response

JSON Preview Visualize

```
{
  "success": true,
  "candidateId": "6a283d5c74ce24f83c05addc",
  "name": "RAJADURAI ARIVAZHAGAN",
  "summary": "The candidate, Rajadurai Arivazhagan, is a strong fit for a Java role given his extensive experience as"
```

Activate Windows
Go to Settings to activate Windows.

Connect Git Console Terminal 0 0 0

Globals Vault Tools

23:26
12-06-2026

File Edit Selection View Go Run Terminal Help

Untitled (Workspace) - Antigravity IDE - README.md

Confirming Vector Embedding Requirements

Continue from where you left

Worked for 57s

I have completed the implementation and verification of the candidate match summarization (Match Explainer) endpoint `/v1/search/summarize`.

Summary of Completed Work

1. AI Service: Implemented the `generateMatchSummary(queryText, candidate)` method in `ai.service.js`:

- Formats the candidate's skills and experience into a text prompt.
- Calls the Groq API completion endpoint (POST `https://api.groq.com/openai/v1/chat/completions`) using the Llama model configured in `.env` (`meta-llama/llama-4-scout-17b-16e-instruct`).
- Extracts and returns the explanation paragraph summarizing why that candidate fits the search query.

2. Controller & Router:

- Implemented `getMatchSummary` in `js/search.controller.js` with validation checks for `q`, presence of `candidateId`, and valid 24-character hexadecimal

0 Files With Changes Review Changes

Ask anything, @ to mention, / for actions

Activate Windows
Go to Settings to activate Windows.

Gemini 3.5 Flash (Medium)

Ln 347, Col 65 (36 selected) Spaces: 2 UTF-8 LF {} Markdown Antigravity - Settings

23:20
12-06-2026

Backend Service > ## API Endpoints Guide > ### 7. Match Summarization API (Match Explainer)

```
1  ## Resume Search Platform: Backend
56  ## API Endpoints Guide
337  ### 7. Match Summarization API (Match Explainer)
382  ---
383
384  ### 8. End-to-End Resume Search Pipeline API
385  * **Endpoints**
386  - `GET /v1/search`
387  - `GET /api/v1/search`
388  * **Query Parameters**
389  - `q` (string, required): The search text/phrase key.
390  * **Use Case**:: Runs RRF hybrid retrieval (combining Vector Search and BM25 Search) to collect relevant candidate profiles, slices the top 5 candidates, triggers a batch LLM explanation summary request using the Groq API (meta-llama/
```

and-88-01-7fkhbk.mongodb.net:27017,ac-9nj1sya-shard-00-02-7fkhbk.mongodb.net:27017/resumedb?ssl=true&replicaSet=atlas-2a9fa6-shard-00&authSource=admin&retryWrites=true&majority=

[2026-06-12 23:17:14.965 +0530] INFO: Server is running in development mode on port 3000



[2026-06-12 23:17:15.407 +0530] INFO: Successfully connected to MongoDB.


rag_resume_retrieval master* 0 0 0

ganeshmynampati (now)

Ln 347, Col 65 (36 selected) Spaces: 2 UTF-8 LF {} Markdown Antigravity - Settings

23:20
12-06-2026





Postman collections for resume retrieval with step by step approaches:

